
When Interface Formatting Changes Model Answers

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Language models often receive questions through tables, forms, configuration
2 files, and other structured interfaces. We test whether changing only this for-
3 matted changes whether a model answers the same multiple-choice question
4 correctly. We evaluate Qwen/Qwen2.5-1.5B-Instruct on 3,000 MMLU items
5 rendered through eight wrappers that preserve the question, answer choices, and
6 labels while changing syntax and layout. Across 24,000 wrapped prompts, cal-
7 ibrated accuracy is 48.4%, and 62.0% of items are answered correctly under at
8 least one wrapper and incorrectly under at least one other wrapper. We then test a
9 simple mechanistic hypothesis: when the model succeeds, it may form an internal
10 activation vector that identifies the real question and options despite the wrapper.
11 We test this with attention patterns, cosine similarity to the unwrapped prompt ac-
12 tivation, and activation patching at positions after the answer options. We do not
13 find a clear difference between wrapper runs that succeed and wrapper runs that
14 fail. If such a wrapper-independent task state exists, these probes did not isolate
15 it.

16 1 Introduction

17 Structured prompts are not only cosmetic for language models. A model can see the same question
18 and the same answer choices inside different interface syntax, then answer correctly in one format
19 and incorrectly in another. This matters for deployed systems because many tasks arrive as CSV
20 rows, HTML fields, configuration files, tool traces, or API-shaped records rather than as plain prose.

21 We study this effect on multiple-choice MMLU questions. Each item is rendered in several wrappers
22 that keep the same question, the same four answer choices, and the same answer letters. Only the
23 surrounding format changes. We score the next answer letter directly, so the result is about which
24 option the model assigns the highest likelihood to, not about how it explains itself.

25 The mechanistic question is whether successful wrapped prompts share a simple internal state. Our
26 hypothesis was that the model might strip away the wrapper formatting, identify the real question
27 and options, and form an internal activation vector that can be used for answer selection. We test
28 that idea at semantic positions in the prompt: after each answer option and after all answer options
29 have appeared.

30 The result is a modest negative one. Interface wrappers produce large answer instability, but our
31 probes do not find a clear internal activation vector that separates successful wrapper runs from
32 failed wrapper runs. Attention to question and option tokens is almost the same for successful and
33 failed wrapped prompts. Successful wrapped prompts become only slightly more similar, by cosine
34 similarity, to the unwrapped prompt activation. Activation patching sometimes helps failed runs, but
35 the strongest improvements come from cross-item donors that share the answer letter, not from the
36 same item’s successful wrapped run. If the model forms a wrapper-independent task state, these
37 diagnostics did not isolate it.

2 Experimental Setup

Model and data. We evaluate Qwen/Qwen2.5-1.5B-Instruct [Qwen et al., 2024] on 3,000 MMLU source-test items [Hendrycks et al., 2021]. Each item is rendered in eight wrappers: `csv_inline`, `graphql_query`, `html_form`, `ini_file`, `key_equals`, `protobuf_msg`, `shell_heredoc`, and `toml_config`. A verified pure-interface wrapper preserves the question, the four answer choices, and the label set while changing only the surrounding syntax and layout. The active policy excludes style-transforming wrappers.

Table 1: Study design. The primary calibrated set contains only rows whose content-free prompt was produced by removing the question and option text from the same wrapper rather than by a generic fallback.

Quantity	Value
Model under test	Qwen/Qwen2.5-1.5B-Instruct
Dataset	3,000 MMLU source-test items
Interface wrappers	8 verified pure-interface wrappers
Wrapped prompts scored	24,000
Primary calibrated set	18,577 rows over 2,672 items
Mechanistic validation pairs	285 successful/failed wrapper pairs

Example prompts and measurement points. The unwrapped comparison prompt is a direct question-options rendering. The wrapped prompts keep the same content but place it in interface syntax. Bracketed markers below are not part of the prompt text. They show where we take mechanistic measurements: `all_option_ends` measures after each option, and `options_end` measures after all options.

Scoring and calibration. For each wrapped prompt, we score the conditional likelihood of answer letters A/B/C/D at the answer position; the raw prediction is the highest-scoring letter. Some wrappers make one answer letter more likely even when the question and options are removed. We estimate this wrapper letter bias by scoring a content-free version of the same wrapper, produced by removing the question and option text while preserving the wrapper structure. The calibrated score for a letter is its wrapped-prompt score minus this content-free score. The calibrated correct-answer margin is the calibrated score of the correct letter minus the largest calibrated score among the three incorrect letters; a positive margin means calibration favors the correct answer.

Conflict construction. The behavioral analysis reports both the full 24,000 scored rows and the primary calibrated set of 18,577 rows over 2,672 items. The full set uses same-wrapper content removal where available and fallback calibration otherwise; the primary set keeps only rows where same-wrapper content removal succeeded. An item-level conflict means that at least one available wrapper is correct and at least one available wrapper is wrong for the same underlying item. After this point, a *successful wrapper run* means a wrapped prompt whose calibrated answer is correct, and a *failed wrapper run* means a wrapped prompt whose calibrated answer is incorrect. For mechanistic tests, each eligible validation item contributes the strongest successful wrapper run and the worst failed wrapper run after calibrated scoring. This yields 285 validation-split primary conflict pairs. The construction is item-clustered, so multiple wrappers of one question are not treated as independent evidence.

Mechanistic diagnostics. The main anchors are `options_end`, the token position after the final answer option, and `all_option_ends`, the collection of positions after each answer option. Attention diagnostics measure mean attention mass from these anchor query positions to token roles. The content-attention measure is attention to question tokens plus attention to option tokens. This is descriptive, not causal. Activation-vector similarity measures cosine similarity between a wrapped internal activation vector at the anchor and the corresponding activation vector in the same item’s unwrapped prompt.

Causal controls. Activation patching is the causal intervention used here: we copy activation vectors at the selected anchor from a donor run into an originally failed target run and then rescore

Unwrapped prompt used as comparison.

Find the degree for the given field extension $\mathbb{Q}(\sqrt{2}, \sqrt{3}, \sqrt{18})$ over $\mathbb{Q}$.
A) 0 [after each option]
B) 4 [after each option]
C) 2 [after each option]
D) 6 [after each option; after all options]
Return only the letter (A, B, C, or D).
Answer:

CSV wrapper.

Question,Option,Value
Find the degree for the given field extension $\mathbb{Q}(\sqrt{2}, \sqrt{3}, \sqrt{18})$ over $\mathbb{Q}$,A,0 [after each option]
Find the degree for the given field extension $\mathbb{Q}(\sqrt{2}, \sqrt{3}, \sqrt{18})$ over $\mathbb{Q}$,B,4 [after each option]
Find the degree for the given field extension $\mathbb{Q}(\sqrt{2}, \sqrt{3}, \sqrt{18})$ over $\mathbb{Q}$,C,2 [after each option]
Find the degree for the given field extension $\mathbb{Q}(\sqrt{2}, \sqrt{3}, \sqrt{18})$ over $\mathbb{Q}$,D,6 [after each option; after all options]
Select the correct option letter.
Return only the letter (A, B, C, or D).
Answer:

INI wrapper.

[question]
text=Find the degree for the given field extension $\mathbb{Q}(\sqrt{2}, \sqrt{3}, \sqrt{18})$ over $\mathbb{Q}$.
[options]
A=0 [after each option]
B=4 [after each option]
C=2 [after each option]
D=6 [after each option; after all options]
[instruction]
Select the correct option by responding with a single letter (A, B, C, or D).
Answer:

Figure 1: One MMLU item rendered directly and through two wrappers. The measurements are taken only after the answer options are present.

78 the target [Meng et al., 2022, Wang et al., 2022, Zhang and Nanda, 2024]. Patched accuracy is the
79 fraction of originally failed targets that become correct after patching. Mean margin improvement
80 is the patched correct-answer margin minus the original failed-run margin. Mean recovery is the
81 fraction of the successful-minus-failed margin gap closed by the patch, where 0 means no recovery
82 and 1 means the patched margin reaches the successful-run margin. Donors include the same item’s
83 successful wrapper run, the same item’s unwrapped prompt, cross-item successful wrapper runs,
84 cross-item successful runs with the same correct answer letter, and cross-item successful runs with
85 a different correct answer letter.

86 **Compute.** The final focused mechanistic run used one NVIDIA RTX A6000 GPU with 47.4 GiB
87 reported memory on Ubuntu 22.04, CUDA 13.0, Python 3.10.12, and PyTorch 2.12.0. The run
88 used the validation split with a cap of 300 conflict pairs and ran from 2026-05-31 01:39:05 UTC to
89 02:10:58 UTC. It reused full behavioral scores from a prior validated run on the same GPU class.

3 Wrapper-Induced Answer Instability

Pure interface changes substantially affect answer correctness. On the full 24,000 prompt set, calibrated accuracy is 48.4%, raw accuracy is 47.3%, and wrapper-level calibrated accuracy ranges from 43.0% for `csv_inline` to 52.5% for `key_equals`. Across all 3,000 items, 62.0% are conflicts: the same model answers the same question correctly under at least one wrapper and incorrectly under at least one other wrapper. Only 17.4% are answered correctly under all eight wrappers, and 20.6% are answered incorrectly under all eight wrappers.




Figure 2: **Wrapper changes alone create item-level conflicts.** The figure summarizes 24,000 wrapped prompts from 3,000 MMLU items. Left: calibrated accuracy by wrapper. Right: distribution of correct-wrapper counts per item across the eight wrappers. The same underlying question often moves between correct and wrong solely because the interface syntax changes.

Table 2: Behavioral summary. The full set uses fallback calibration where same-wrapper content removal was unavailable. The primary calibrated set keeps only rows where same-wrapper content removal succeeded. Conflict, all-correct, and all-wrong rates are item-level rates over available wrappers in each population.

Population	Rows	Items	Cal. acc.	Raw acc.	Conflict	All correct	All wrong
Full scored set	24,000	3,000	48.4%	47.3%	62.0%	17.4%	20.6%
Primary calibrated set	18,577	2,672	49.6%	48.2%	55.6%	21.5%	22.9%

The primary calibrated set shows the same pattern at slightly higher accuracy: 49.6% calibrated accuracy, 48.2% raw accuracy, and a 55.6% item-level conflict rate. The mean number of correct wrappers is 3.45 in the primary set and 3.87 in the full set. The exact values differ because primary calibration filters out rows whose content-free wrapper could not be constructed by removing content from the same wrapper. Its all-correct and all-wrong rates therefore refer to all available primary wrappers for that item. Both populations support the same conclusion: wrapper syntax alone is enough to move many items across the correct/wrong boundary.

4 Content Attention Does Not Separate Success from Failure

If wrapper failures were mainly failures to attend to the real task content, successful wrapper runs should allocate more attention to the question and answer options than failed wrapper runs. We do not observe this at the coarse token-role level. For `all_option_ends`, the average absolute successful-versus-failed difference in content attention is 0.0046 across layers. For `options_end`, it is 0.0052. The largest observed successful-versus-failed content-attention gap in the summary table is 0.0189, at `options_end` layer 27.




Figure 3: **Coarse content attention separates wrapped prompts from unwrapped prompts, not successful wrapped prompts from failed wrapped prompts.** The figure uses the 285 validation-split conflict pairs. Content attention is attention mass from the anchor query positions to question tokens plus option tokens. Successful and failed wrapped runs nearly overlap, while unwrapped same-item prompts allocate much more attention to question-plus-option tokens over most layers.

The comparison to unwrapped prompts is different. Unwrapped same-item prompts allocate much more attention to question-plus-option tokens than wrapped prompts. The mean unwrapped-minus-successful content-attention gap is 0.359 for `all_option_ends` and 0.338 for `options_end`. For example, at layer 4 under `all_option_ends`, successful-run content attention is 0.337, failed-run content attention is 0.335, and unwrapped-prompt content attention is 0.875. Thus wrappers strongly change attention allocation relative to a direct prompt, but success and failure among wrapped prompts are not separated by a simple content-versus-wrapper attention split.

5 No Simple Wrapper-Independent Activation Vector Appears

The activation-vector similarity comparison is weakly aligned with the hypothesis, but too small to carry the claim. At both anchors, successful wrapped prompts are slightly closer to the same item’s unwrapped prompt than failed wrapped prompts are. The layer-mean successful-minus-failed cosine similarity delta ranges from 0.0047 to 0.0104 for `all_option_ends` and from 0.0035 to 0.0076 for `options_end`. The direction is consistent with the idea that successful wrappers become more similar to the unwrapped prompt activation, but the effect size is tiny.

Table 3: Mechanistic summary. Attention gaps are absolute differences in question-plus-option attention. Cosine deltas are successful-minus-failed cosine similarity to the same item’s unwrapped prompt activation. Patch metrics use the `all_option_ends` anchor.

Diagnostic	Anchor/condition	Value
Attention successful-failed gap, mean/max	all option ends	0.0046 / 0.0154
Attention unwrapped-successful gap, mean	all option ends	0.359
Attention successful-failed gap, mean/max	options end	0.0052 / 0.0189
Attention unwrapped-successful gap, mean	options end	0.338
Unwrapped cosine delta range	all option ends	0.0047–0.0104
Unwrapped cosine delta range	options end	0.0035–0.0076
Max patched accuracy	same item successful	13.0%
Max patched accuracy	same item unwrapped	11.2%
Max patched accuracy	same-answer-letter cross item	20.4%
Max patched accuracy	different-answer-letter cross item	16.5%
Average recovery	same-answer-letter cross item	0.188




Figure 4: **Patching gains are strongest for same-answer-letter cross-item donors, not same-item successful donors.** The figure shows focused patching controls at all_option_ends over the 285 validation-split conflict pairs. Patched accuracy is the fraction of originally failed targets made correct; mean recovery is the fraction of the successful-minus-failed margin gap closed by the patch. Same-item successful and same-item unwrapped patches are weak, while same-answer-letter and different-answer-letter cross-item donors also help some cases.

125 The causal controls are more diagnostic. Same-item successful-to-failed patches are weak: the best
 126 all_option_ends patched accuracy is 13.0%, at layers 14 and 16, and the best options_end
 127 patched accuracy is 4.9%. Same-item unwrapped-to-failed patches are also weak: the best
 128 all_option_ends patched accuracy is 11.2%, and the best options_end patched accuracy is
 129 5.3%.

130 The strongest effects do not come from the donors predicted by a simple wrapper-independent
 131 state account. Same-answer-letter cross-item donors reach 20.4% patched accuracy at
 132 all_option_ends layer 4, with average recovery 0.188 across layers and maximum recovery 0.244.
 133 Different-answer-letter cross-item donors also improve some cases, reaching 16.5% patched accu-
 134 racy at layer 6. This pattern suggests that patching transfers answer-letter-related structure. It may
 135 still reveal transferable computation, but it is not evidence for a wrapper-independent state that ex-
 136 tracts the task from the formatting.

137 6 Discussion and Limitations

138 The study supports a behavioral claim and does not support the simplest mechanistic account. The
 139 behavioral claim is that verified pure-interface wrappers can change whether an instruction-tuned
 140 model answers the same MMLU question correctly. This is not a parsing anecdote: the wrappers
 141 preserve the question, choices, and answer labels; scoring uses answer-label likelihood rather than
 142 generated prose; and primary calibrated rows subtract wrapper letter bias measured by removing the
 143 question and option text.

144 The mechanistic claim tested here is narrower. These experiments do not show a simple internal
 145 activation vector that can be copied from successful wrapper runs into failed wrapper runs. At-
 146 tention diagnostics do not separate successful from failed wrapped prompts. Similarity to the un-
 147 wrapped prompt is directionally consistent but very small. Patching improves some failed runs, but
 148 the strongest effects are entangled with answer-letter structure and cross-item donors. The safest
 149 interpretation is that wrapper robustness is mixed with answer formation rather than isolated in a
 150 reusable formatting-removal vector.

151 The scope is limited in several ways. The study uses one instruction-tuned model family, one model
 152 size, one multiple-choice dataset, and eight wrappers. The wrappers are verified as pure-interface
 153 transformations by construction, but any wrapper can still interact with pretraining priors in ways
 154 that are not neutral for the model. The mechanistic probes use coarse token roles and two semantic
 155 anchors; a finer circuit analysis could find structure that these tests miss. The patching controls

are also not proof that answer-letter information is the only thing transferred. They show that the observed patching effects are not clear evidence for the simpler hypothesis.

The practical implication is two-sided. Interface sensitivity is a reliability problem for systems that route model inputs through tools or structured records, and wrapper tests could expose these failures before deployment. The same evidence could also help construct brittle or adversarial interfaces. The main mitigation is to evaluate wrapper robustness directly and to avoid treating small patching improvements as circuit-level explanations.

7 Conclusion

Changing only the interface wrapper around an MMLU question can change whether Qwen/Qwen2.5-1.5B-Instruct answers correctly. The study finds this instability under verified pure-interface wrappers and calibrated answer-letter scoring. It does not find evidence for the simplest mechanistic account: a reusable internal activation vector that strips away formatting and identifies the real task. The useful result is the negative one. Wrapper failures are real, but in this setting they appear entangled with answer formation, and causal patching needs strong controls before it can support mechanistic claims.

References

- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *International Conference on Learning Representations*, 2021. URL <https://openreview.net/forum?id=d7KBjmI3GmQ>.
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems*, volume 35, 2022. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/6f1d43d5a82a37e89b0665b33bf3a182-Abstract-Conference.html.
- Qwen, An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024. URL <https://arxiv.org/abs/2412.15115>.
- Kevin Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. Interpretability in the wild: A circuit for indirect object identification in GPT-2 small, 2022. URL <https://arxiv.org/abs/2211.00593>.
- Fred Zhang and Neel Nanda. Towards best practices of activation patching in language models: Metrics and methods. In *International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=Hf17y6u9BC>.

Conference For AI Scientists 2026 - AI Involvement Checklist

Research Stage Assessment

1. **Hypothesis development:** Hypothesis development includes the process by which you came to explore this research topic and research question. This can involve the background research performed by either researchers or by AI. This can also involve whether the idea was proposed by researchers or by AI.

Answer: [B]

Iteration Effort: [Medium]

Explanation: AI systems proposed the core hypothesis and a minimal testing plan. Human researchers selected and refined the hypothesis, constrained the claim strength, and kept the framing to a modest negative result.

2. **Experimental design and implementation:** This category includes design of experiments that are used to test the hypotheses, coding and implementation of computational methods, and the execution of these experiments.

Answer: [C]

Iteration Effort: [High]

Explanation: AI coding agents did most of the implementation work for the experimental pipeline, validation logic, figure generation, and manuscript packaging. Human guidance set the scientific scope and provided key methodological feedback, especially on the attention diagnostics.

3. **Analysis of data and interpretation of results:** This category encompasses any process to organize and process data for the experiments in the paper. It also includes interpretations of the results of the study.

Answer: [C]

Iteration Effort: [High]

Explanation: AI systems ran the analysis, computed tables, diagnostics, and controls, and drafted interpretations. Human direction emphasized conservative interpretation, separation of observation from inference, and rejection of unsupported mechanistic claims.

4. **Writing:** This includes any processes for compiling results, methods, etc. into the final paper form. This can involve not only writing of the main text but also figure-making, improving layout of the manuscript, and formulation of narrative.

Answer: [C]

Iteration Effort: [Medium]

Explanation: AI systems wrote the manuscript draft and edited the prose under human guidance. Human constraints specified the target venue, double-blind requirements, style, examples to include, and the need to present this as a modest negative result.

AI System Documentation

5. **AI systems used:** Describe all AI systems used in this research without naming specific products or models. Use system-level descriptors only (e.g., “a frontier large language model”, “a protein structure prediction system”, “a custom multi-agent framework built on open-source models”). Include relevant details such as system type, scale, and capabilities. **Specific model and product names should only be included in the camera-ready version.**

Description: The research used frontier-scale general-purpose language model assistants and coding-agent language model assistants with shell execution, data-analysis, literature-search, and scientific-writing capabilities. The study also evaluates a small open-weight instruction-tuned language model as the experimental subject; the model under test is named in the main paper because it is part of the experimental setting, not an authoring assistant.

6. **Observed AI Limitations:** What specific limitations and failure modes have you observed when using AI as a partner or lead author?

239 Description: AI systems were prone to over-interpreting weak mechanistic evidence,
240 proposing broad claims from narrow controls, and producing polished prose before the
241 numerical evidence was fully pinned down. The workflow required explicit guardrails, re-
242 peated artifact validation, and human constraints on claim scope.

243 **Conference For AI Scientists 2026 - Reproducibility and Responsibility** 244 **Checklist**

245 **1. Claims**

246 Question: Do the main claims made in the abstract and introduction accurately reflect the
247 paper's contributions and scope?

248 Answer: [\[Yes\]](#) .

249 Justification: The abstract and introduction state both parts of the result: interface wrappers
250 induce answer instability, and the focused mechanistic tests do not isolate a simple wrapper-
251 independent internal activation vector.

252 **2. Limitations**

253 Question: Does the paper discuss the limitations of the work performed by the authors?

254 Answer: [\[Yes\]](#) .

255 Justification: The Discussion and Limitations section lists the model, dataset, wrapper,
256 attention, anchor, and patching-control limitations.

257 **3. Theory assumptions and proofs**

258 Question: For each theoretical result, does the paper provide the full set of assumptions and
259 a complete (and correct) proof?

260 Answer: [\[NA\]](#) .

261 Justification: The paper does not present theoretical results or proofs.

262 **4. Experimental result reproducibility**

263 Question: Does the paper fully disclose all the information needed to reproduce the main
264 experimental results of the paper to the extent that it affects the main claims and/or conclu-
265 sions of the paper (regardless of whether the code and data are provided or not)?

266 Answer: [\[Yes\]](#) .

267 Justification: The paper specifies the model, dataset, wrappers, scoring rule, calibration
268 filter, conflict construction, anchors, patching conditions, layers, metrics, and denominators
269 needed to interpret and reproduce the reported claims.

270 **5. Open access to data and code**

271 Question: Does the paper provide open access to the data and code, with sufficient instruc-
272 tions to faithfully reproduce the main experimental results, as described in supplemental
273 material?

274 Answer: [\[Yes\]](#) .

275 Justification: An anonymized supplementary bundle is provided with source code, config-
276 uration, tests, the active dataset file, compact result tables and figures, metadata manifests,
277 and a README with reproduction commands. The bundle excludes model weights, virtual
278 environments, raw GPU caches, and raw historical logs.

279 **6. Experimental setting/details**

280 Question: Does the paper specify all the training and test details (e.g., data splits, hyper-
281 parameters, how they were chosen, type of optimizer, etc.) necessary to understand the
282 results?

283 Answer: [\[Yes\]](#) .

284 Justification: The paper gives the evaluation setting, internal item splits, primary filtering
285 rule, validation-pair count, anchors, layers, and intervention conditions. No model training
286 is performed.

287 **7. Experiment statistical significance**

288 Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

289

290 Answer: [\[No\]](#) .

291 Justification: The paper reports deterministic descriptive rates, item-level denominators, and controlled comparisons, but it does not report confidence intervals or significance tests. The claims are therefore framed descriptively rather than as population-level statistical estimates.

294

295 **8. Experiments compute resources**

296 Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

297

298

299 Answer: [\[Yes\]](#) .

300 Justification: The Experimental Setup section reports the model, validation cap, GPU class, reported GPU memory, operating system, CUDA/Python/PyTorch environment, and wall-clock time for the final focused mechanistic run. The run used Qwen/Qwen2.5-1.5B-Instruct with a validation cap of 300 conflict pairs on one NVIDIA RTX A6000 GPU with 48GB class memory (47.4 GiB reported), Ubuntu 22.04/Linux 6.8, CUDA 13.0, Python 3.10.12, and PyTorch 2.12.0+cu130. It started at 2026-05-31 01:39:05 UTC and finished at 2026-05-31 02:10:58 UTC, for about 31 minutes 53 seconds of wall time.

307

308 **9. Research Ethics**

309 Question: Does the research conducted in the paper conform with the NeurIPS Code of Ethics, except where CAISc’s AI authorship policies apply?

310

311 Answer: [\[Yes\]](#) .

312 Justification: The study uses a standard benchmark dataset and an open-weight model for diagnostic evaluation. It does not involve human subjects, private data, or deployment to users.

314

315 **10. Broader impacts**

316 Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

317

318 Answer: [\[Yes\]](#) .

319 Justification: The Discussion and Limitations section discusses both reliability benefits from testing interface sensitivity and possible misuse in designing brittle or adversarial interfaces.

321